

4F03 Assignment 1
Parallel Image Processing

Jamie Counsell
Rakesh Mistry

Thursday, Jan 28, 2015

4F03 Assignment 1

This project is also available through [GitHub](#). The URL will be private until the submission deadline.

Authors

Name	Student Number	Email	Website
Rakesh Mistry	1221428	mistryrn@mcmaster.ca	rakeshmistry.ca
Jamie Counsell	1054209	mistryrn@mcmaster.ca	jamiecounsell.me

Description

This C program provides two methods of image filtering to reduce noise in `ppm` formatted images. The two filter options are:

- [Mean Filter](#)
- [Median Filter](#)

Dependencies

- `openmpi`

Installation

- Install the dependencies
- Clone the repository

Operation

To run the program, first run make:

```
$ make clean
... [truncated]
$ make
mpicc -g -O2 -Wall -Wno-unused-variable -c -o main.o main.c
mpicc -g -O2 -Wall -Wno-unused-variable -c -o readwriteppm.o readwriteppm.c
mpicc -g -O2 -Wall -Wno-unused-variable -c -o processimage.o processimage.c
mpicc -o ppmf main.o readwriteppm.o processimage.o
```

Then run the program as follows:

```
$ mpirun -np p ./ppmf input.ppm output.ppm N F
```

where:

- **p** is the number of processes
- **input.ppm** is the name of the input file
- **output.ppm** is the name of the output file
- **N** specifies the size of the window, that is $N \times N$ window, where N is an odd integer ≥ 3
- **F** is the type of filter. **A** for *mean filter*, and **M** for *median filter*

Work Breakdown

The work was broken down as follows:

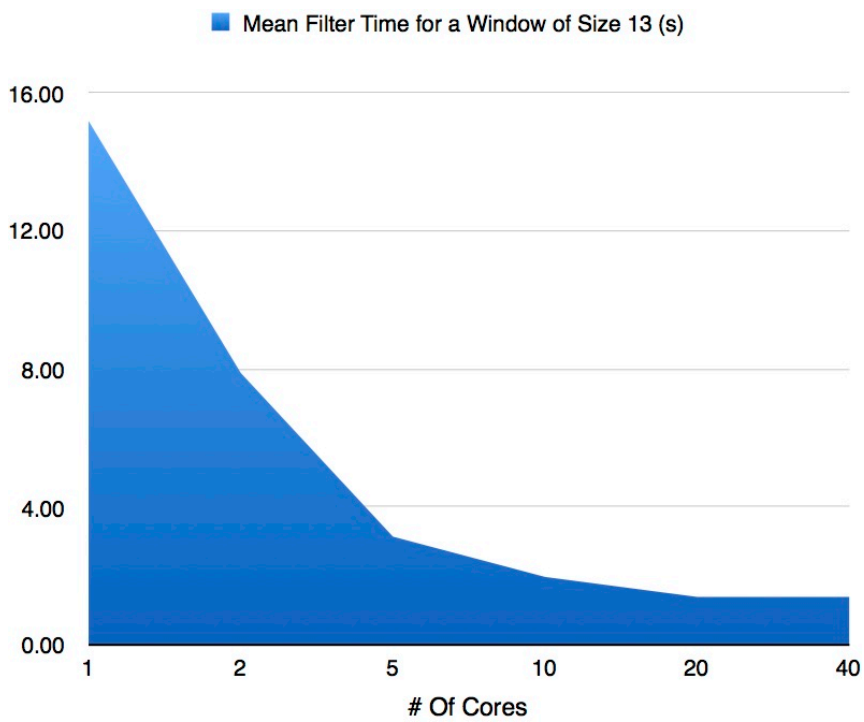

- Jamie wrote the mean filter
- Rakesh wrote the median filter
- Both authors wrote the rest of `main.c` and `processimage.c` together, in person

Notes

- Each process allocates and receives only the chunk of the image that it requires. Process 0 allocates and stores the entire source image, as it needs it to write back at the end of the filtering. Each other process uses only the pixels required to compute its ration of the original problem.
- This current implementation will suffer if an image loaded is larger than available memory. It is assumed that this is not the case for this assignment.

Results

The following are results from McMaster's MPIHost server.



main.c

```
#include "a1.h"
#include "mpi.h"
#include <time.h>
#include <stdlib.h>
#include <stdio.h>
#include <assert.h>

int main(int argc, char** argv)
{
    // Initialize variables
    RGB *image = NULL;
    int width = 0, height = 0, max = 0;
    int my_rank, p, i, start, size;
    clock_t begin = 0;
    double time_spent;
    int window = atoi(argv[3]);

    // Initialize MPI
    MPI_Status status;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &p);

    if (my_rank == 0) {
        // Rank 0 read image from disk
        image = readPPM(argv[1], &width, &height, &max);
    }

    // Broadcast width and height to other processes
    MPI_Bcast(&width, 1, MPI_INT, 0, MPI_COMM_WORLD);
    MPI_Bcast(&height, 1, MPI_INT, 0, MPI_COMM_WORLD);

    if (my_rank == 0) {
        begin = clock();
        for (i=1; i<p; i++) {
            start = getStart(i, width, height, window, p);
            size = getSize(i, width, height, window, p);
            MPI_Send(image + start, size*sizeof(RGB), MPI_CHAR, i, 0, MPI_COMM_WORLD);
        }
    } else {
        start = getStart(my_rank, width, height, window, p);
        size = getSize(my_rank, width, height, window, p);
        image = (RGB*)malloc(size*sizeof(RGB));
        assert(image);
        MPI_Recv(image, size*sizeof(RGB), MPI_CHAR, 0, 0, MPI_COMM_WORLD, &status);
    }
    processImage(width, height, image, argc, argv);

    if (my_rank == 0) {
        time_spent = (double)(clock() - begin) / CLOCKS_PER_SEC;
        writePPM(argv[2], width, height, max, image);
        printf("done. processImage took %fs\n", time_spent);
    }

    // Cleanup
    free(image);
    MPI_Finalize();

    return(0);
}
```

processImage.c

```
#include "a1.h"
#include "mpi.h"
#include <stdio.h>
#include <stdlib.h>

void processImage(int width, int height, RGB *image, int argc, char** argv)
{
    // Initialize variables
    int i, my_rank, p, n, start, my_size, process_start, process_size;
    double a, b;

    // Initialize MPI
    MPI_Status status;
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &p);
    int source = 0, dest = 0, tag = 0;

    // Get total size of image
    int size = width*height;

    // Get number of pixels per process
    double h = size/p;

    // Get parameters from command line
    int window = atoi(argv[3]);
    char *filter = argv[4];

    // Print header
    if (my_rank == 0) {
        printf("Window size:      %dx%d\nFilter Type:      %s\n", window, window, filter);
        printf("Pixels in image: %d\nPixels/process:  %.0f\nTotal processes: %d\n", size, h, p);
        printf("-----\n");
    }

    /* create a type for struct RGB */
    const int nitems=3;
    int blocklengths[3] = {sizeof(char), sizeof(char), sizeof(char)};
    MPI_Datatype types[3] = {MPI_CHAR, MPI_CHAR, MPI_CHAR};
    MPI_Datatype mpi_rgb_type;
    MPI_Aint offsets[3];
    offsets[0] = offsetof(RGB, r);
    offsets[1] = offsetof(RGB, g);
    offsets[2] = offsetof(RGB, b);
    MPI_Type_create_struct(nitems, blocklengths, offsets, types, &mpi_rgb_type);
    MPI_Type_commit(&mpi_rgb_type);
    /* ----- */

    MPI_Barrier(MPI_COMM_WORLD);
    // Determine lower and upper values for pixel range
    start = getStart(my_rank, width, height, window, p);
    my_size = getSize(my_rank, width, height, window, p);
    if (my_rank == 0) {
        process_start = 0;
        process_size = my_size - (width * (window-1)/2);
    } else if (my_rank >= p-1) {
        process_start = (width * (window-1)/2);
        process_size = my_size;
    } else {
        process_start = (width * (window-1)/2);
        process_size = my_size;
    }
}
```

processImage.c

```
}

// Run filtering on image
printf("Process %d crunching pixels %d - %d...\n", my_rank, start, my_size);

if ( *filter == 'A' ) { // mean filter
    meanFilter(my_size, width, image, window, process_start, process_size, my_rank);

} else if ( *filter == 'M' ) { // median filter
    medianFilter(my_size, height, image, window, process_start, process_size, my_rank);

} else { // Invalid input for filter type
    if (my_rank == 0){
        printf("Error: Invalid filter specified. Please use either 'A' for Mean, or 'M' for
Median.\n");
    }
}

if (my_rank != 0) {
    // Send this rank's image chunk to process zero
    MPI_Send(image + process_start, process_size, mpi_rgb_type, dest, tag, MPI_COMM_WORLD);
} else {
    for (i=1; i < p; i++) {
        start = size/p*i;
        process_size = getSize(i, width, height, window, p);
        // Wait to receive data from other processes
        MPI_Recv(image + start, process_size, mpi_rgb_type, i, tag, MPI_COMM_WORLD, &status);
    }
}
}
```

median.c

```
#include "a1.h"
#include <stdlib.h>
#include <string.h>
#include <stdint.h>
#include <stdio.h>

float median(int n, int x[]);
float median2(int n, int arr[]);

void swap(int* a, int* b);

void medianFilter(int size, int width, RGB *image, int window, int start, int end, int rank){
    int thing = (window - 1)/2;
    int windowsq = window*window;
    int topleft, current;
    int pc, i, j, cypx;

    // Storage for the R, G, and B values in the window
    int rvalues[windowsq];
    int gvalues[windowsq];
    int bvalues[windowsq];

    RGB *unmodified = (RGB*)malloc(size*sizeof(RGB));
    RGB *current_pixel;
    RGB *copydestpixel;
    RGB *copysrcpixel;
    RGB *pixel;

    // Deep copy unmodified <- img
    for (i=0; i < size; i++) {
        copydestpixel = unmodified + i;
        copysrcpixel = image + i;

        copydestpixel->r = copysrcpixel->r;
        copydestpixel->g = copysrcpixel->g;
        copydestpixel->b = copysrcpixel->b;
    }

    // For each pixel in this worker's quota
    for (pc = start; pc < end; pc++) {
        // Current pixel of interest
        pixel = image + pc;

        // Pixel at top left of window
        topleft = pc - (thing * width) - thing;

        // Reset R, G, and B values to 0's
        memset(rvalues, 0, sizeof(int)*windowsq);
        memset(gvalues, 0, sizeof(int)*windowsq);
        memset(bvalues, 0, sizeof(int)*windowsq);

        int count = 1;

        // For each row in window
        for (i=0; i < window; i++) {
            // For each column in window
            for (j=0; j < window; j++) {
                // Determine the pixel we're looking at
                current = topleft + i * width - thing + j + 1;
```


median.c

```
// If current pixel is outside range of window and image, skip it
if (current < 0 || current > size - 1 || (current % width) > (pc % width) + thing) {
    // Do nothing

    // If current pixel is in range of window and image
} else {
    current_pixel = unmodified + current;
    rvalues[count - 1] = current_pixel->r;
    gvalues[count - 1] = current_pixel->g;
    bvalues[count - 1] = current_pixel->b;

    count = count + 1;
}
}
}

pixel->r = median2(count, rvalues);
pixel->g = median2(count, gvalues);
pixel->b = median2(count, bvalues);
}
}

// Code courtesy of https://en.wikiversity.org/wiki/C\_Source\_Code/Find\_the\_median\_and\_mean
// Great for testing, but slow.
float median(int n, int x[]) {
    float temp;
    int i, j;
    // the following two loops sort the array x in ascending order
    for (i = 0; i < n-1; i++) {
        for (j = i+1; j < n; j++) {
            if (x[j] < x[i]) {
                // swap elements
                temp = x[i];
                x[i] = x[j];
                x[j] = temp;
            }
        }
    }

    if (n%2 == 0) {
        // if there is an even number of elements, return the mean of the two elements in the
        middle
        return ((x[n/2] + x[n/2 - 1]) / 2.0);
    } else {
        // else return the element in the middle
        return x[n/2];
    }
}

void swap(int* a, int* b){
    int t = *a;
    *a = *b;
    *b = t;
}

// Code courtesy of http://stackoverflow.com/questions/1961173/median-function-in-c-math-library
// Much faster than median().
float median2(int n, int arr[])
{
    int low, high ;
```

median.c

```
int median;
int middle, ll, hh;
low = 0 ; high = n-1 ; median = (low + high) / 2;
for (;;) {
    if (high <= low) /* One element only */
        return arr[median] ;
    if (high == low + 1) { /* Two elements only */
        if (arr[low] > arr[high])
            swap(arr+low, arr+high) ;
        return arr[median] ;
    }
    /* Find median of low, middle and high items; swap into position low */
    middle = (low + high) / 2;
    if (arr[middle] > arr[high]){
        swap(arr + middle, arr +high);
    }
    if (arr[low] > arr[high]){
        swap(arr+low, arr+high);
    }
    if (arr[middle] > arr[low]){
        swap(arr+middle, arr+low);
    }

    /* Swap low item (now in position middle) into position (low+1) */
    swap(arr+middle, arr+low+1) ;
    /* Nibble from each end towards middle, swapping items when stuck */
    ll = low + 1;
    hh = high;

    for (;;) {
        do ll++; while (arr[low] > arr[ll]) ;
        do hh--; while (arr[hh] > arr[low]) ;
        if (hh < ll)
            break;
        swap(arr+ll, arr+hh) ;
    }

    /* Swap middle item (in position low) back into correct position */
    swap(arr+low, arr+hh) ;
    /* Re-set active partition */
    if (hh <= median)
        low = ll;
    if (hh >= median)
        high = hh - 1;
}
return arr[median];
}
```

mean.c

```
#include "a1.h"
#include <stdlib.h>

void meanFilter(int size, int width, RGB *image, int window, int start, int end, int rank){
    int thing = (window - 1)/2;
    int topleft, current;
    int pc, i, j, cypx;
    double sum[3];
    RGB *unmodified = (RGB*)malloc(size*sizeof(RGB));
    RGB *current_pixel;
    RGB *copydestpixel;
    RGB *copysrcpixel;
    RGB *pixel;

    // Deep copy to store unmodified values
    for (i=0; i < size; i++) {
        copydestpixel = unmodified + i;
        copysrcpixel = image + i;
        copydestpixel->r = copysrcpixel->r; // SEG FAULT HERE
        copydestpixel->g = copysrcpixel->g;
        copydestpixel->b = copysrcpixel->b;
    }

    // For each pixel in this worker's quota
    for (pc = start; pc < end; pc++) {
        // Current pixel of interest
        pixel = image + pc;

        // Pixel at top left of window
        topleft = pc - (thing * width) - thing;

        sum[0] = 0; // Red values
        sum[1] = 0; // Green values
        sum[2] = 0; // Blue values
        int count = 0;
        //printf("sum: %0.1f   count: %d\n", sum, count);

        // For each row in window
        for (i=0; i < window; i++) {
            // For each column in window
            for (j=0; j < window; j++) {
                // Determine the pixel we're looking at
                current = topleft + i * width - thing + j + 1;

                // If current pixel is outside range of window and image, skip it
                if (current < 0 || current > size - 1 || (current % width) > (pc % width) + thing) {
                    // Do nothing

                    // If current pixel is in range of window and image
                } else {
                    current_pixel = unmodified + current;
                    sum[0] = sum[0] + (current_pixel->r);
                    sum[1] = sum[1] + (current_pixel->g);
                    sum[2] = sum[2] + (current_pixel->b);

                    count = count + 1;
                }
            }
        }
    }
}
```

mean.c

```
//printf("pc: %d    total:    %f\n", pc, sum);  
// new pixel rgb value is average of ^  
pixel->r = sum[0]/count;  
pixel->g = sum[1]/count;  
pixel->b = sum[2]/count;  
}  
}
```

helpers.c

```
// Get start position of chunk for process
int getStart(int rank, int width, int height, int window, int p) {
    int total = width * height;
    int ration = total/p;

    int start = ration * rank - ((window - 1)/2 * width);

    if (start < 0) {
        start = 0;
    }
    return start;
}

// Get size of chunk for process
int getSize(int rank, int width, int height, int window, int p) {
    int total = width * height;
    int ration = total/p;
    int size = ration + (width * (window-1));
    int start = getStart(rank, width, height, window, p);
    if (rank >= p-1 || size > total) {
        size = total - start;
    } else if (start == 0) {
        size = size - ((window-1)/2 * width);
    }
    return size;
}
```

readwriteppm.c

```
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>
#include "mpi.h"
#include "al.h"

void writePPM(char* file, int width, int height, int max, const RGB *image)
{
    int i;
    printf("Writing result... ");

    // open file for writing
    FILE *fd;
    fd = fopen(file, "w");

    // output the header
    fprintf(fd, "P3\n");
    fprintf(fd, "%d %d\n%d\n", width, height, max);
    // write the image
    for (i = 0; i < height*width; i++)
    {
        const RGB *p = image+i;
        fprintf(fd, "%d %d %d ", p->r, p->g, p->b);
    }
}

RGB * readPPM(char* file, int* width, int* height, int* max)
{
    /* Read a PPM P3 image from a file into a buffer. Return the
       buffer, width, height, and the max value in the image. */

    FILE *fd;
    char c;
    int i, n, my_rank;

    // initialize MPI
    MPI_Status status;
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
    int source, total, dest = 0, tag = 0;

    char b[100];
    int red, green, blue;

    // check if P3 ppm format
    fd = fopen(file, "r");
    n = fscanf(fd, "%[^\n] ", b);
    if (b[0] != 'P' || b[1] != '3')
    {
        if (my_rank == 0) {
            printf("%s is not a PPM file!\n", file);
        }
        exit(0);
    }
    if (my_rank == 0) {
        // printf("%s is a PPM file\n", file);
    }
    n = fscanf(fd, "%c", &c);
    while (c == '#')
    {
        n = fscanf(fd, "%[^\n] ", b);
    }
}
```

readwriteppm.c

```
    if (my_rank == 0) {
        printf("%s\n", b);
    }
    n = fscanf(fd, "%c", &c);
}
ungetc(c, fd);
n = fscanf(fd, "%d %d %d", width, height, max);
assert(n==3);

// size of image
int size = *width*(*height);
RGB *image = (RGB*)malloc(size*sizeof(RGB));
assert(image);

for(i=0; i < size; i++)
{
    n = fscanf(fd, "%d %d %d", &red, &green, &blue);
    assert(n==3);
    image[i].r = red;
    image[i].g = green;
    image[i].b = blue;
}

return image;
}
```

a1.h

```
#ifndef INCLUDED_A1_H
#define INCLUDED_A1_H

typedef struct {
    unsigned char r,g,b;
} RGB;

RGB * readPPM(char* file, int* width, int* height, int* max);
void writePPM(char* file, int width, int height, int max, const RGB *image);

void meanFilter(int size, int width, RGB *image, int window, int start, int end, int rank);
void medianFilter(int size, int width, RGB *image, int window, int start, int end, int rank);

int getSize(int rank, int width, int height, int window, int p);
int getStart(int rank, int width, int height, int window, int p);

void processImage(int width, int height, RGB *image, int argc, char** argv);

#endif
```


makefile

```
CFLAGS=-g -O2 -Wall -Wno-unused-variable  
CC=mpicc
```

```
PROGRAM_NAME= ppmf  
OBJS = main.o readwriteppm.o processimage.o mean.o median.o helpers.o
```

```
$(PROGRAM_NAME): $(OBJS)  
    $(CC) -o $@ $?
```

```
clean:  
    rm *.o $(PROGRAM_NAME) *~
```